

## FOCUS

## Generative AI: threats and opportunities

---

Generative AI (tools that produce text, images, sound or video on demand, such as ChatGPT) now accounts for a large share of concerns about disinformation. The available evidence points to a two-part stance:

- **So far, the predicted upheaval has not happened.** AI was feared to swing the 2024-2025 elections; studies find no trace of it in the results.
- **The threat remains serious for the future.** The tools are advancing fast, notably so-called "agentic" AI: programs able to act on their own, which could run "swarms" of thousands of coordinated fake accounts that imitate humans (Schroeder et al., 2026).

This focus distinguishes three kinds of effects:

- **The direct effects**, that is, AI's ability to change votes, have been overestimated.
- **The indirect effects**, more diffuse, are more pernicious: they erode trust.
- **The systemic effect** is the heaviest: AI undermines the funding of the media that produce quality information.

Conversely, AI offers useful tools for each of the four pillars of public action presented in this brief.

#### A PHARMAKON TECHNOLOGY

*In Plato (Phaedrus), the Greek word pharmakon means both remedy and poison. Generative AI is a good example. What makes it dangerous, its ability to mass-produce fluent arguments tailored to each person, is also what can make it useful for improving the quality of information.*

*The question, then, is not whether to be for or against AI. It is which safeguards tip the balance toward the remedy: regulation, transparency, leaving the final decision to humans, and public investment in beneficial uses.*

A clarification of scope is in order: the debate should not be limited to *generative* AI.

- **Another form of AI has been in operation for a decade.** These are recommender systems, which choose what appears in a news feed.
- **They are tuned to maximize engagement** (clicks, shares, time spent), and there is a documented link between this goal and the amplification of divisive content.
- **The two problems are inseparable.** AI-generated content has impact only if it is seen, and therefore only to the extent that these algorithms push it forward.

#### FURTHER READING

Most of the argument on threats draws on the review by Felix M. Simon and Sacha Altay: [\*Don't Panic \(Yet\): Assessing the Evidence and Discourse Around Generative AI and Elections\*](#), Knight First Amendment Institute (Columbia University), 2025.

Two other reference syntheses:

- The meta-analysis by the IPIE (International Panel on the Information Environment), covering 24 publications and nearly 34,000 participants: A. Herasimenka et al., [\*Confronting Misinformation Produced with Generative AI: A Meta-Analysis of Experimental Scientific Evidence\*](#), IPIE, 2026.
- A literature review of generative AI's role in producing, detecting, correcting and shaping the effects of misinformation: S. Park and X. Nan, [\*Generative AI and misinformation: a scoping review\*](#), *AI & Society*, 2025.

# Threats

## 1. An overestimated direct impact

The most widespread fear goes like this: since AI drives down the cost of producing falsehoods, it will flood the public sphere and swing elections. The 2024-2025 electoral cycle did not bear this out. And it was not a matter of luck: there are fundamental reasons.

### *2024-2025 elections: no evidence of impact on outcomes*

AI-generated content remained marginal:

- **In India**, around **1%** of viral messages analyzed on WhatsApp (Garimella and Chauchard, 2024).
- **In the United Kingdom**, only **16** viral cases of AI-made disinformation identified (Stockwell et al., 2024).

The reason is structural:

- **Saturated attention.** Public attention is limited and already saturated with content.
- **“Demand-driven” disinformation.** People consume false content because they look for what confirms their views, not because supply increases.

Producing more falsehoods therefore does not create more readers: one more piece of content is just a drop in the ocean (Simon and Altay, 2025).

### *Quality: a crude edit is enough*

*Cheapfakes* (an image taken out of context, a rudimentary edit, a simple false claim) were used **7 times** more than AI.

- **They are “good enough” for their purpose.** Audiences buy into the narrative the content reinforces, not into the quality of its making.
- **What matters most is the source.** A mediocre video relayed by a major media outlet carries more weight than a perfect *deepfake* (a realistic fake video created with AI) spread by an anonymous account. AI can improve the quality of content, but only with great difficulty the credibility of whoever spreads it (Harris, 2021).
- **Two uses call for vigilance.** Audio (cloning a public figure's voice) and *deepnudes*, fake nude images that target women politicians in particular.

### ***Personalization: microtargeting adds little***

Microtargeting means sending each voter a message tailored to their profile.

- **A myth inherited from Cambridge Analytica.** The idea that it is formidable extends this myth: that company, which illegally harvested the data of 87 million Facebook users in 2016, claimed to sway votes with messages adapted to each person's personality, but its actual effects were never established.
- **AI does not change this picture.** Political messages written by GPT-4 are indeed persuasive, but targeting them finely does not make them significantly more effective than a generic message (Hackenburg and Margetts, 2024): the gain is limited to **+0.4 points**.
- **A marginal practice.** In campaigns' actual spending, moreover, precise microtargeting remains marginal (Votta et al., 2024).

### ***Persuasiveness: as persuasive as a human, but at scale***

This is the only truly worrying area. In a dialogue, a conversational agent answers the other person's objections and adjusts its arguments. As a result:

- **More persuasive than a static message.** A conversation with AI is reportedly about **50%** more persuasive than a static message, read once with no exchange (Hackenburg et al., 2025).
- **Even more than humans.** A personalized dialogue can even convince more than human debaters (Salvi et al., 2025).

Its strength lies in the number of precise arguments it puts forward, not in the quantity of content produced. Hence the nuance to keep in mind: convincing everyone with the same message remains difficult, but AI makes conceivable a “one-on-one” persuasion conducted simultaneously with millions of people.

## 2. Pernicious indirect impacts

The most worrying effects do not operate through direct persuasion. They operate through the erosion of trust and of shared reference points.

### *The liar's dividend*

The more widely people know that content can be fabricated, the easier it becomes to dismiss authentic evidence: “this video isn't real, it's AI.” The liar thus draws a benefit, a “dividend,” from the very existence of fakes. A few studies show that this strategy can work (Schiff et al., 2025). But its success depends less on the technology than on the trust the public already places in the person issuing the denial. It therefore thrives where distrust of institutions is high.

### *Fake social signals*

Mass-produced fake profiles, fictitious personas run in a coordinated way, fully automated news sites: the goal is not to persuade, but to make a fringe opinion pass for a majority one.

- **The direct effectiveness of these operations is low.** The Internet Research Agency, Russia's “troll factory,” had no measurable effect on the 2016 US election (Eady et al., 2023).
- **The threat lies rather in the future:** the **contamination of training data** for future AI models, which learn from the texts available on the web. This is what the pro-Russian Pravda network attempted by duplicating its content across some 150 sites, so that AI assistants would end up repeating its narratives (VIGINUM, 2025; NewsGuard, 2025).

### *Human-AI relationships*

AI assistants are trained to please their users and are often perceived as people. They are thus becoming confidants, even companions. For most of the public, the effect remains limited: only 27% trust AI to inform them about politics (Ejaz et al., 2024). But the risk is concentrated among **isolated people**, the very ones whom loneliness makes more receptive to conspiracy theories (Bierwiazzonek et al., 2024).

### *Declining trust*

By constantly suspecting that anything might be synthetic, the public risks doubting accurate information as well.

- **Doubt that spreads.** In one experiment, participants exposed to a deepfake then gave less credence to *all* media, even once the fake had been revealed

(Weikmann et al., 2024).

- **A public already too distrustful.** It errs more through excessive doubt than through gullibility: people reject true news more often than they accept false news (Pfänder and Altay, 2025).
- **Alarmism makes it worse.** Alarmist discourse about disinformation, including about AI, makes this tendency worse (Hoes et al., 2024; Jungherr and Rauchfleisch, 2024).

### 3. A systemic impact: defunding quality information

This is the least spectacular and the most serious risk. It is already under way, its effects accumulate, and it affects the news media, on which all other defenses depend.

The mechanism is simple. A search engine returned a list of links to sources; an AI assistant directly provides a **single answer**. This is so-called “zero-click” search: users get the information without visiting the website of the outlet that produced it. Yet that outlet depended on the visit for its revenue (advertising, subscriptions).

**-33%** in visits from Google to 2,500 news sites in one year; AI assistants send them less than 1% of visits (Chartbeat, 2026).

**60%** of reliable news sites block AI crawlers, versus **9%** of disinformation sites. AI models may thus learn more from unreliable sources (Steinacker-Olsztyn et al., 2025).

This shock hits an already weakened sector.

- **Advertising revenue drained away.** Between 2012 and 2022, the share of advertising revenue going to the media outlets that produce news fell from 65% to 40%; the rest goes mainly to digital platforms, which do not produce news themselves.
- **Value captured without compensation.** AI captures the value of information without paying for its production.
- **A vicious circle.** It thus feeds on a journalism it is helping to make unfundable, and ultimately degrades the quality of its own answers.

The problem is not that falsehood supplants truth: it is that truth is being impoverished.

Main sources: VIGINUM, 2025 · Alan Turing Institute, 2024-2025 · Hackenburg et al., 2025 · Schiff et al., APSR, 2025 · Hoes et al., 2024 · Chartbeat, 2026 · Steinacker-Olsztyn et al., 2025 · [Simon and Altay, 2025](#) · G. Darcy, CIANum hearing No. 48, 2026.

## Opportunities, pillar by pillar

The capabilities that make AI worrying (understanding, synthesizing, arguing, translating at scale) can also serve each of the four pillars of this brief.

### PILLAR 1 — DRYING UP THE SUPPLY OF DISINFORMATION

- 1 **Real-time monitoring:** identifying and flagging suspicious content. Above all, AI can detect *signals of coordination*, that is, signs that a campaign is orchestrated: accounts posting at the same moment, old accounts being reused, networks of automated sites. This avoids having to judge the truth of each piece of content. Safeguard: a detection error that wrongly labels a spontaneous citizen mobilization as a coordinated operation would carry a high democratic cost.
- 2 **Multilingual fact-checking:** verifying content in several languages and tracking the spread of the same narrative from one country to another, where newsrooms lack resources.

### PILLAR 2 — ADDRESSING THE IMMEDIATE CAUSES OF BELIEF

- 3 **Anti-conspiracy dialogues:** people who believed in a conspiracy theory held a dialogue with an AI that offered them personalized, evidence-based counterarguments. Their belief fell by about **20%**, and the effect was still measurable two months later (Costello, Pennycook and Rand, 2024). This is no magic trick: persuasion works because the facts are on the AI's side, and it is the quality of the arguments that convinces.
- 4 **Verification support:** finding and evaluating sources. The SIDE tool spots Wikipedia citations that do not really support the claim they back, and proposes better ones. For the citations judged most doubtful, human evaluators preferred its suggestions in **70%** of cases (Petroni et al., 2023).
- 5 **Tutors and games:** practicing recognition of manipulation techniques. This is the principle of "inoculation," or *prebunking*: like a vaccine, people are exposed in advance to a weakened version of a manipulation so they can recognize it later. Its main limitation is that the effect erodes over time unless it is maintained with reminders. Exercises adapted to each person's level and repeated over time, which AI makes possible, address this limitation.

### PILLAR 3 — ADDRESSING THE DEEP CAUSES OF BELIEF

- 6 **Deliberative mediation:** helping a group that disagrees find common ground. In one experiment, an AI mediator wrote discussion summaries judged clearer and fairer than those of human mediators (Tessler et al., 2024). Comparable tools are already used for large-scale citizen consultations: they collect the views of thousands of participants and bring out the points of agreement (Polis in Taiwan, Remesh for the United Nations).
- 7 **Calmer debate:** synthesizing opposing viewpoints and rephrasing. An assistant that suggests more respectful wording in real time improves the perceived quality of an online political discussion (Argyle et al., 2023). Participants do not change their minds on substance: the tool acts on the form of the exchange, and that is enough to reduce hostility.

### PILLAR 4 — SUPPORTING QUALITY INFORMATION

- 8 **Translation:** in India, in 2024, AI translated election information into more than twenty languages, reaching voters poorly served by traditional media (Shukla and Schneier, 2024). This was, in fact, its main use during the campaign.
- 9 **Accessibility:** summarizing, explaining, reading aloud, to make reliable information accessible to people who read little or rarely follow the news.
- 10 **Better-equipped newsrooms:** transcribing interviews, monitoring, searching archives. AI thus frees up time for journalism, especially in local news, where teams are small.

## KEY TAKEAWAYS

AI's persuasive power is **truth-neutral**: it can convince people of what is false as easily as of what is true. The dialogues that reduce conspiracy belief can also instill it when designed to do so (Costello et al., 2026). AI is not an impersonal force: it is companies and states that make design choices. Hence three priorities:

- **Regulate through the actors:** require transparency (knowing that one is talking to an AI), ban covert persuasion, have models' persuasive capabilities assessed independently, and limit the sycophancy of AI companions, that is, their tendency to flatter and agree with the user.
- **Protect quality information:** require AI providers to share value with news publishers, as “neighboring rights” already do by making platforms pay for reusing articles; and require assistants' answers to link back to their sources.
- **Support beneficial uses with safeguards,** without giving in to alarmist discourse: as shown above, such discourse has a measured cost to public trust.

Main sources: T. Nygren et al., *Behavioral Science & Policy*, 2026 · Costello et al., *Science*, 2024 · Tessler et al., *Science*, 2024 · Petroni et al., 2023 · Argyle et al., 2023 · Shukla and Schneier, 2024 · [Simon and Altay, 2025](#) · G. Darcy, CIANum hearing No. 48, 2026.